

# 6.036 Notes Spring 2024

## ① types of learning

supervised

- regression

- classification

structural - bad model

error

estimation - lack of data

error

unsupervised

- clustering

- density estimation (probability distribution)

- dimensionality reduction

sequence = mapping from input sequence to output sequence

reinforcement = maps states to optimal actions based on rewards

## Loss

$$0-1 \text{ Loss} = \begin{cases} 0 & g=a \\ 1 & \text{guess } g \neq \text{answer } a \end{cases}$$

$$\text{Squared Loss} = (g-a)^2$$

$$\text{Absolute Loss} = |g-a|$$

$$\text{Asymmetric Loss} = \begin{cases} x & \text{if } g=1 \text{ and } a=0 \\ y & \text{if } g=0 \text{ and } a=1 \\ 0 & \text{else} \end{cases}$$

to prioritize some mistakes over others

usually we want to minimize expected loss

## Matrix Derivatives

$\frac{\partial y}{\partial x}$  is  $1 \times 1$  scalar

$\frac{\partial y}{\partial x}$  is  $\overset{n \times 1}{\begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \vdots \\ \frac{\partial y}{\partial x_n} \end{bmatrix}}$   $n \times 1$  vector

$\overset{m \times 1}{\frac{\partial y}{\partial x}}$  is  $\begin{bmatrix} \frac{\partial y_1}{\partial x} & \frac{\partial y_2}{\partial x} & \dots & \frac{\partial y_m}{\partial x} \end{bmatrix}$   $1 \times m$  vector

$\overset{n \times m}{\frac{\partial y}{\partial x}}$  is  $\begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_2}{\partial x_1} & \dots & \frac{\partial y_m}{\partial x_1} \\ \vdots & & & \vdots \\ \frac{\partial y_1}{\partial x_n} & \frac{\partial y_2}{\partial x_n} & \dots & \frac{\partial y_m}{\partial x_n} \end{bmatrix}$   $n \times m$  matrix

$$\frac{\partial \mathbf{x}}{\partial \mathbf{A}}_{n \times m} := \begin{bmatrix} \frac{\partial \mathbf{x}}{\partial A_{11}} & \frac{\partial \mathbf{x}}{\partial A_{12}} & \cdots & \frac{\partial \mathbf{x}}{\partial A_{1m}} \\ \vdots & \ddots & \ddots & \vdots \\ \frac{\partial \mathbf{x}}{\partial A_{n1}} & \ddots & \ddots & \frac{\partial \mathbf{x}}{\partial A_{nm}} \end{bmatrix}_{n \times m} \text{ matrix}$$

$$\frac{\partial \mathbf{a}^T}{\partial \mathbf{x}} = 0 \quad \xrightarrow{\text{unrelated to } \mathbf{x}}$$

$$\frac{\partial \mathbf{x}}{\partial \mathbf{x}} = I$$

$$\frac{\partial \mathbf{A}\mathbf{x}}{\partial \mathbf{x}} = \mathbf{A}^T$$

$$\frac{\partial \mathbf{A}^T \mathbf{A}}{\partial \mathbf{x}} = \mathbf{A}$$

$$\frac{\partial \mathbf{v} \mathbf{u}}{\partial \mathbf{x}} = \mathbf{v} \frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}} \mathbf{u}^T \quad \leftarrow \text{Product Rule.}$$

$$\frac{\partial g(\mathbf{u})}{\partial \mathbf{x}} = \frac{\partial g(\mathbf{u})}{\partial \mathbf{u}} \cdot \frac{\partial \mathbf{u}}{\partial \mathbf{x}} \quad \leftarrow \text{Chain Rule}$$

$$\frac{\partial \mathbf{u}^T \mathbf{v}}{\partial \mathbf{x}} = \frac{\partial \mathbf{u}}{\partial \mathbf{x}} \mathbf{v} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}} \mathbf{u} \quad \frac{\partial \mathbf{u}^T}{\partial \mathbf{x}} = \left( \frac{\partial \mathbf{u}}{\partial \mathbf{x}} \right)^T$$

## ② Linear regression

incorporate  $\theta_0$  into here via extra coordinate

$$J(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n (\boldsymbol{\theta}^T \tilde{\mathbf{x}}^{(i)} - y^{(i)})^2$$

$$\nabla_{\boldsymbol{\theta}} J = \frac{2}{n} \tilde{\mathbf{X}}^T (\tilde{\mathbf{X}} \boldsymbol{\theta} - \tilde{\mathbf{Y}}) \Rightarrow \text{set to 0}$$

$$\Rightarrow \boxed{\boldsymbol{\theta} = (\tilde{\mathbf{X}}^T \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^T \tilde{\mathbf{Y}}}$$

$$\tilde{\mathbf{X}} = n \begin{bmatrix} \text{one point} \\ \vdots \\ \text{all first coord.} \end{bmatrix}$$

$$\tilde{\mathbf{Y}} = n \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix} \text{ actual}$$

$$\boldsymbol{\theta} = \begin{bmatrix} \theta_0 \\ \vdots \\ \theta_d \end{bmatrix} \text{ prediction parameters}$$

regularize by ridge regression

$$\Rightarrow \boldsymbol{\theta} = (\tilde{\mathbf{X}}^T \tilde{\mathbf{X}} + n \lambda \mathbf{I})^{-1} \tilde{\mathbf{X}}^T \tilde{\mathbf{Y}}$$

for

$$J_{\text{ridge}}(\boldsymbol{\theta}, \boldsymbol{\theta}_0) = \frac{1}{n} \sum_{i=1}^n (\boldsymbol{\theta}^T \tilde{\mathbf{x}}^{(i)} - y^{(i)})^2 + \lambda \|\boldsymbol{\theta}\|^2$$

$\Rightarrow$  avoids singularities.

## Cross Validation

- break  $D$  into  $k$  batches  $D_k$ 
  - for each  $k$ , train using  $D/D_k$  and validate on  $D_k$ .

return  $\frac{1}{k} \sum_{i=1}^k [\text{error using model trained with } D/D_k \text{ on } D_k]$



## Gradient Descent

more in direction opposed to gradient,  $\nabla f$  at a time.

when not improving more than  $\epsilon$  per step, **STOP**.

→ many dimensions.

$$\nabla_{\theta} f = \begin{bmatrix} \partial f / \partial \theta_1 \\ \vdots \\ \partial f / \partial \theta_m \end{bmatrix}$$

$$\theta^{(t)} = \theta^{(t-1)} - \eta \nabla_{\theta} f(\theta^{(t-1)})$$

**STOP** at  $|f(\theta^{(t)}) - f(\theta^{(t-1)})| < \epsilon$ ,

## for linear regression

$$\theta^{(t+1)} = \theta^{(t-1)} - \eta \frac{2}{n} \sum_{i=1}^n \left( [\theta^{(t-1)}]^T x^{(i)} - y^{(i)} \right) x^{(i)}$$

$$\sum \nabla_{\theta} J$$

## for ridge regression

$$J_{\text{ridge}}(\theta, \theta_0) = \frac{1}{n} \sum_{i=1}^n (\theta^T x^{(i)} + \theta_0 - y^{(i)})^2 + \lambda \|\theta\|^2$$

$$\nabla_{\theta} J_{\text{ridge}}(\theta, \theta_0) = \frac{2}{n} \sum_{i=1}^n (\theta^T x^{(i)} + \theta_0 - y^{(i)}) x^{(i)} + 2\lambda \theta$$

update  
 $\theta, \theta_0$   
separately

$$\frac{\partial J_{\text{ridge}}(\theta, \theta_0)}{\partial \theta_0} = \frac{2}{n} \sum_{i=1}^n (\theta^T x^{(i)} + \theta_0 - y^{(i)})$$

## stochastic gradient descent

$$f(\theta) = \sum_{i=1}^{\infty} f_i(\theta)$$

objective

mini-batch

pick a random parameter  $i$  to change

$$f_i \Rightarrow \theta^{(t)} = \theta^{(t-1)} - \eta(t) \nabla_{\theta} f_i(\theta^{(t-1)})$$

random sample  
size  $\lceil \frac{n}{k} \rceil$  from  $n$  data points

if  $f$  convex,  $\eta(t)$  satisfies

$$\sum_{t=1}^{\infty} \eta(t) = \infty$$

$$\sum_{t=1}^{\infty} \eta^2(t) < \infty,$$

stochastic gradient descent converges with probability 1 to optimal  $\theta$ .

④ linear classifier  $\rightarrow$  chooses 1 or 0,  
 $E_n(h) = \frac{1}{n} \sum_{i=1}^n \begin{cases} 1 & h(x^{(i)}) \neq y^{(i)} \\ 0 & \text{otherwise} \end{cases}$

training  
error

usually,  $h(x; \theta, \theta_0) = \text{sign}(\theta^T x + \theta_0)$ ,

$\hookrightarrow$  but usually NP-hard to find  $\theta, \theta_0$ .

sigmoid  $\sigma(z) = \frac{1}{1+e^{-z}}$

can predict 1 for  $\sigma(z) > 0.5$ , 0 otherwise

where  $z = \theta^T x + \theta_0$

want to maximize

$$\prod_{i=1}^n \left\{ \begin{array}{ll} g^{(i)} & \text{if } y^{(i)} = 1 \\ 1-g^{(i)} & \text{otherwise} \end{array} \right. \rightarrow \left. \begin{array}{l} \text{P(born 1)} \\ \text{P(born 0)} \end{array} \right\} \text{right}$$

$$= \prod_{i=1}^n g^{(i)}^{y^{(i)}} (1-g^{(i)})^{1-y^{(i)}}$$

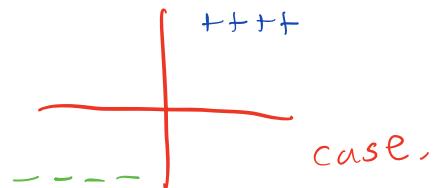
maximized  
when

$$\log \left( \prod_{i=1}^n g^{(i)}^{y^{(i)}} (1-g^{(i)})^{1-y^{(i)}} \right) \text{ is maximized}$$

i.e. minimize  $\left[ \sum_{i=1}^n -y^{(i)} \log g^{(i)} - (1-y^{(i)}) \log (1-g^{(i)}) \right] = \sum_{i=1}^n L_{\text{neg}}(g^{(i)}, y^{(i)})$

must have  
regularization

→ otherwise  $\theta \rightarrow \infty$  is  
optimal in



$$J_{\text{er}}(\theta, \theta_0; D) = \left( \frac{1}{n} \sum_{i=1}^n L_{\text{neg}}(\sigma(\theta^T x^{(i)} + \theta_0), y^{(i)}) \right) + \lambda \|\theta\|^2$$

$$\nabla_{\theta} J_{\text{er}}(\theta, \theta_0) = \frac{1}{n} \sum_{i=1}^n (g^{(i)} - y^{(i)}) x^{(i)} + 2\lambda \theta$$

$$\frac{\partial J_{\text{er}}(\theta, \theta_0)}{\partial \theta_0} = \frac{1}{n} \sum_{i=1}^n (g^{(i)} - y^{(i)})$$

↳ gradient descent!! possible since NLL (negative log likelihood)  
function is convex.

multiple classes ( $K$ , instead of 2)

- training label = one-hot vector = all 0's but 1 at  $i^{\text{th}}$  spot  
if in class  $i$ .

- finding  $\theta, \theta_0$  s.t.  $z = \theta^T x + \theta_0$  then

$$y = h(x; \theta, \theta_0) = \text{softmax}(z) = \begin{bmatrix} \exp(z_1) / \sum \exp(z_i) \\ \vdots \\ \exp(z_K) / \sum \exp(z_i) \end{bmatrix}$$

$$L_{\text{negm}}(g, y) = - \sum_{i=1}^K y_i \cdot \log(g_i)$$

↑  
probability distribution vector

when evaluating accuracy, we check

$$A(h; D) = 1 - \frac{1}{n} \sum_{i=1}^n L_{0,1}(g^{(i)}, y^{(i)})$$

no probability  
deterministic

## ⑤ Feature Representation

$[x] \rightarrow [x, x^2, x^3]$  ← gives more possibility for separability.

$[x, y, z] \rightarrow [x^2, xy, y^2, yz, z^2, xz, x, y, z]$

can also encode using dataset  $D$  with  $n$  points  $x^{(1)}, x^{(2)}, \dots, x^{(n)}$

$$\phi([x_1, x_2, \dots, x_d]) \mapsto \left[ f_{x^{(1)}}(x), f_{x^{(2)}}(x), \dots, f_{x^{(n)}}(x) \right]^T, f_p(x) = e^{-\beta \|p - x\|^2}$$

discrete encodings

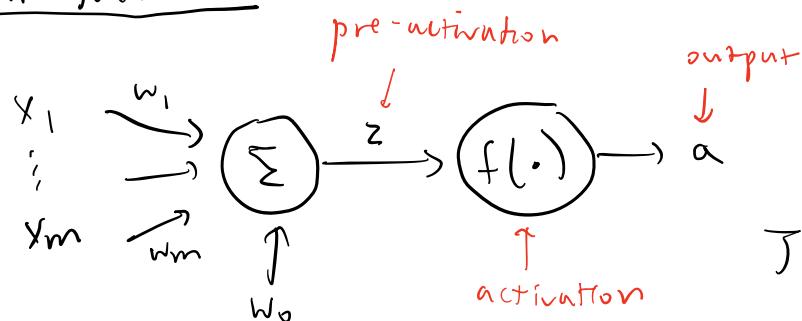
- ordering + mapping → numeric encoding
- ordering but no natural mapping → thermometer
- factoring ex) model + year → factor into two features, encode each one
- no ordering / categorical → one-hot
- binary code  $(00, 01, 10, 11)$  is usually BAD.

normalization  
can help:  $\phi(x) = \frac{x - \bar{x}}{\sigma}$

$(0, 0, 0, 0)$   
 $(1, 0, 0, 0)$   
 $(1, 1, 0, 0)$   
 $(1, 1, 1, 0)$   
 $(1, 1, 1, 1)$

text → bag of words  
(one-hot for each word)

## ⑥ Neural Networks



for one neuron  
want to minimize

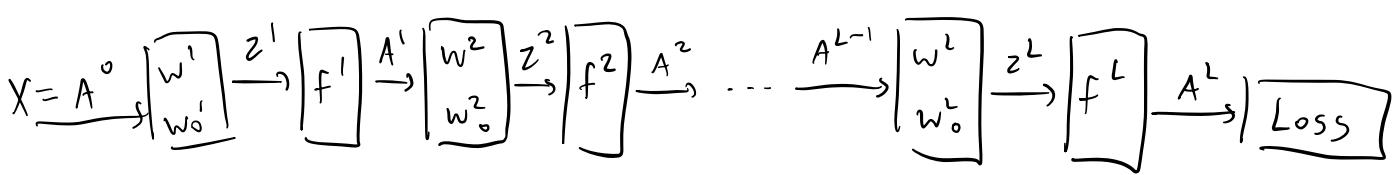
$$J(w, w_0) = \sum_i L(NN(x^{(i)}, w, w_0), y^{(i)})$$

many neurons per layer

using  $A = f(Z) = f(W^T X + w_0)$

activation  $(m \times n)^T$  matrix  $n \times 1$  vector

# inputs to  $l^{th}$  layer =  $m^l$   
# outputs to  $l^{th}$  layer =  $n^l = m^{l+1}$



## activation

$$\text{step}(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases} \quad \text{ReLU}(z) = \begin{cases} 0 & \text{if } z < 0 \\ z & \text{if } z \geq 0 \end{cases} \quad g(z) \quad \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} \quad \text{Softmax}(z)$$

## back-propagation

$$\frac{\partial \text{loss}}{\partial w^l} = A^{l-1} \left( \frac{\partial \text{loss}}{\partial z^l} \right)^T$$

$m^l \times n^l \quad m^l \times 1 \quad 1 \times n^l$

$$\begin{aligned} \frac{\partial \text{loss}}{\partial z^l} &= \frac{\partial A^l}{\partial z^l} \cdot \frac{\partial z^{l+1}}{\partial A^l} \cdot \frac{\partial A^{l+1}}{\partial z^{l+1}} \cdots \frac{\partial A^L}{\partial z^L} \cdot \frac{\partial \text{loss}}{\partial A^L} \\ &= \frac{\partial A^l}{\partial z^l} \cdot W^{l+1} \cdot \frac{\partial A^{l+1}}{\partial z^{l+1}} \cdots W^L \cdot \frac{\partial \text{loss}}{\partial A^L} \end{aligned}$$

do forward pass, then  
calculate "blame"

$$\frac{\partial \text{loss}}{\partial A^l} = \begin{cases} \frac{\partial z^{l+1}}{\partial A^l} \cdot \frac{\partial \text{loss}}{\partial z^{l+1}} & \text{if } l < L \\ \frac{\partial \text{loss}}{\partial A^L} & \text{else} \end{cases}$$

$$\frac{\partial \text{loss}}{\partial z^l} = \frac{\partial A^l}{\partial z^l} \cdot \frac{\partial \text{loss}}{\partial A^l}$$

$$\begin{aligned} \frac{\partial \text{loss}}{\partial w^l} &= A^{l-1} \cdot \left( \frac{\partial \text{loss}}{\partial z^l} \right)^T \\ \frac{\partial \text{loss}}{\partial w_0^l} &= \frac{\partial \text{loss}}{\partial z^l} \end{aligned}$$

$$w^l = w^l - \eta(f) \frac{\partial \text{loss}}{\partial w^l}$$

$$w_0^l = w_0^l - \eta(f) \frac{\partial \text{loss}}{\partial w_0^l}$$

## parameter optimization

- batches
- adaptive step-size

- regularization
- early stopping
- ridge regression
- negative feedback
  - on  $\|W\|^2$
- add random noise to training data
- normalize data

- dropout
  - for each  $j$ , consider  $a_j^l$ 's for each  $l$  and set to 0 with probability  $p$  ( $= \frac{1}{2}$  usual) per training iteration

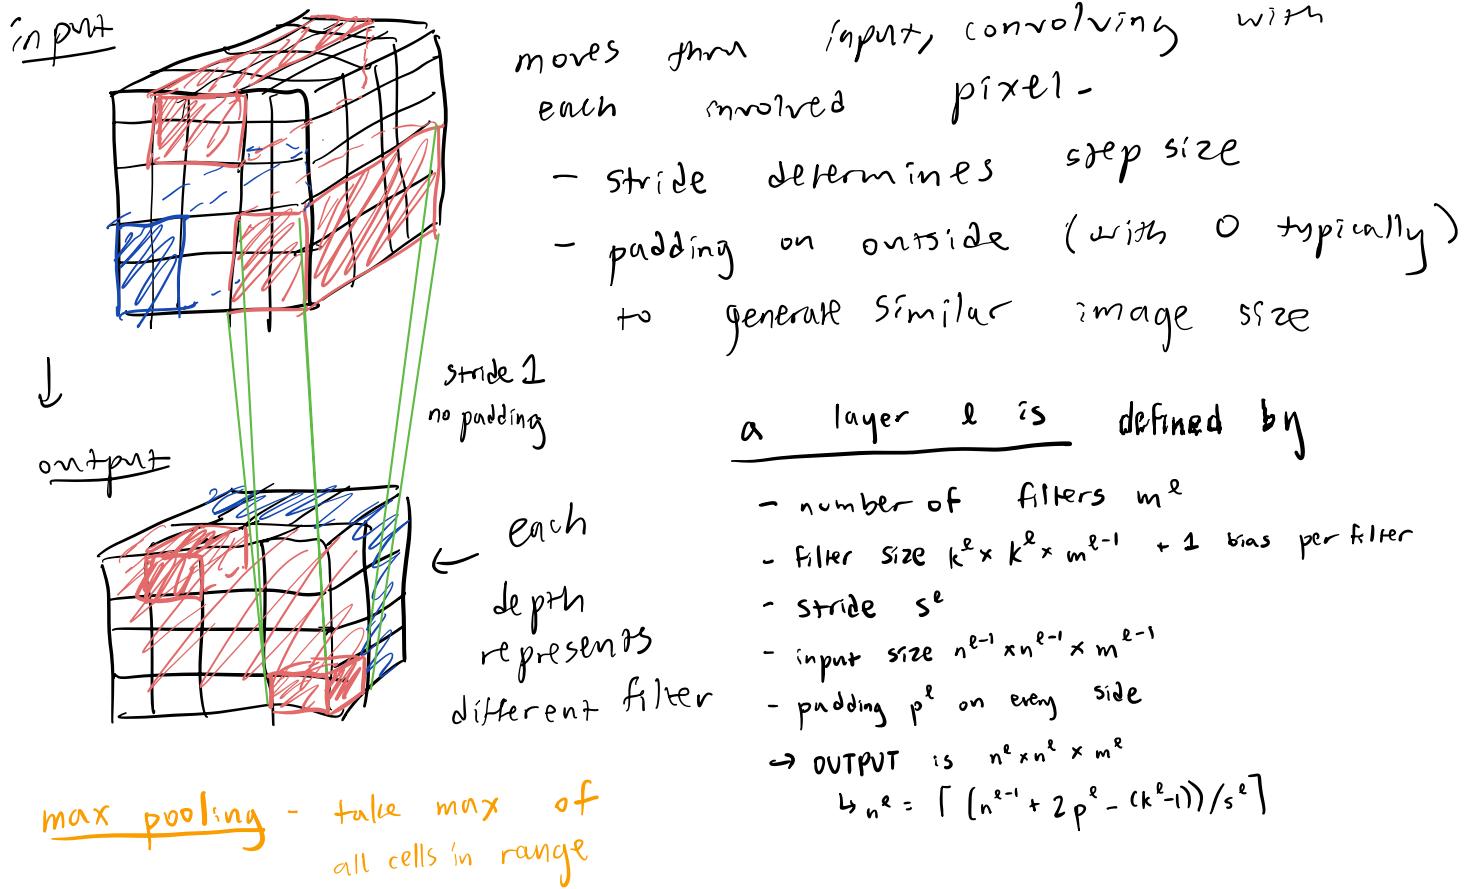
## ⑦ Convolutional Neural Networks

spatial locality - convolutional layer

translational invariance - max pooling layer

### Filters

Starting with  $n \times n$  image, separate into  $n \times n \times 3$  based on RGB.



usually: input  $\rightarrow$  (convolution + ReLU + pooling)  $\rightarrow$  Fully-connected  $\rightarrow$  softmax many times

### backpropagation

for something like  $Z_i^l = W^l A_{[i-(k/2):i+(k/2)]},$

$i^{\text{th}}$  column of  $\frac{\partial Z^l}{\partial W^l} = A_{[i-(k/2):i+(k/2)]} \xleftarrow{\text{column vector}}^{k \times 1}$

$$\nabla_{(x,y)} \max(x, y) = \begin{cases} (1, 0) & x > y \\ (0, 1) & y > x \end{cases}$$

## ⑧ Transformers

semantic vector embeddings :  $v_{paris} - v_{france} \approx v_{rome} - v_{italy}$   
word2vec

$$p(k|q) = \text{softmax}([q^T k_1, q^T k_2, \dots, q^T k_n])$$

↑                      ↑  
 query                other  
 keys

→ roughly higher softmarks  
for closer keys

$$\text{attention}(q, k, v) = \sum_k p(k|q) v(k)$$

↑      ↑      ↑  
 query   key   value      ↑      ↑  
 all other      respective  
 keys      value

mask can be finite which tokens are used in computation

Formally,  $\text{L}$  transformer blocks

transformer block:  $\mathbb{R}^{n \times d} \rightarrow \mathbb{R}^{n \times d}$

params       $n = \# \text{ tokens}$        $d = \text{token dimension}$

$\theta$  = model parameters

each token  $x^{(i)} \in \mathbb{R}^{d \times 1}$

Q( $x^{(i)}$ )    K( $x^{(i)}$ )    V( $x^{(i)}$ )  
 query              key              value  
 J                      P                      ↑

learnable weight matrices

$$W_{h,q}, \quad W_{h,k}, \quad W_{h,r} \in \mathbb{R}^{d \times d_k}$$

$h$  is an index over each transformer head ( $1, 2, \dots, H$ )

$W_{h,c} \in \mathbb{R}^{d_k \times d}$  weight matrix for heads

$$W_1 \in \mathbb{R}^{d \times m} \quad W_2 \in \mathbb{R}^{m \times d}$$

## ① attention weights

$$\alpha_{ij} = \text{softmax}_j \left( \frac{\mathbf{Q}^T(\mathbf{x}^{(i)}) \mathbf{K}(\mathbf{x}^{(j)})}{\sqrt{d_K}} \right)$$

$$\textcircled{2} \quad y^{(l_i)} = \sum_{j=1}^n \alpha_j^{(h)} V(x^{(s)})$$

$$(3) \quad u^{(l,i)} = \sum_{h=1}^H W_{h,c}^T y^{(l)}$$

$$\textcircled{4} \quad u^{(i)} = \text{LayerNorm}(x^{(i)} + u^{(i)}; \gamma_i, \beta_i)$$

$$(5) \quad z^{(l_i)} = w_2^T \text{ReLU}(w_1^T u^{(l)})$$

$$(6) \quad \underline{\text{Final}} \quad z^{(i)} = \text{LayerNorm} \left( z^{(c_i)} + u^{(c_i)} ; \gamma_2, \beta_2 \right)$$

## Non Parametric Methods

$k$ -nearest neighbor: classify each point in region as majority vote of  $k$  nearest training points

Decision trees: focus on CART/ID3 - axis-aligned and fit constant model in leaf

## ① regression

$$I_m := \{ i \mid x^{(i)} \in R^m \}$$

split up into  $m$  regions  $R_1, R_2, \dots, R_m$

each region has output  $o_1, o_2, \dots, o_m \rightarrow$  classifier

$$\hookrightarrow o_m = \text{average}_{i \in I_m} y^{(i)}$$

$\uparrow$   
actual values.

$$\hookrightarrow \text{Error } E_m = \sum_{i \in I_m} (y^{(i)} - o_m)^2$$

$$\text{minimize } \sum_{i=1}^m E_i + \lambda m \leftarrow \begin{array}{l} \text{avoid} \\ \text{many regions} \end{array}$$

algorithm: keep breaking regions into smaller regions until all are  $\leq k$  in size. divide region along single line minimizing

$$\sum_{i \in I_1} (y^{(i)} - y_1^{\text{avg}})^2 + \sum_{i \in I_2} (y^{(i)} - y_2^{\text{avg}})^2.$$

if tree is too large, "weakest-link" pruning to remove the bottom-level split minimizing  $C_\alpha(T) = \sum_{m=1}^{|T|} E_m(T) + \alpha |T|$

## ② classification

$$o_m = \text{majority}_{i \in I_m} y^{(i)} \quad E_m = \# \{ i \mid i \in I_m \text{ and } y^{(i)} \neq o_m \}$$

$$\text{empirical probability } \hat{P}_{m,k} = 1 - \frac{E_m}{N_m} \leftarrow |I_m|$$

### possible metrics

misclassification error

$$Q_m(T) = \frac{E_m}{N_m}$$

Gini index

$$Q_m(T) = \sum_k \hat{P}_{m,k} (1 - \hat{P}_{m,k})$$

$\underbrace{\quad}_{\text{over all classes } k}$

Entropy

$$Q_m(T) = H(I_m) = - \sum_k \hat{P}_{m,k} \log_2 \hat{P}_{m,k}$$

minimize by maximizing information gain

$$H(I_m) = \left( \frac{|I_1|}{N_m} \cdot H(I_1) + \frac{|I_2|}{N_m} \cdot H(I_2) \right)$$

$\uparrow$   
fixed.

## Bootstrap Aggregating (Bagging)

make multiple trees w/small samples,  
decide as average of trees' outputs

## random forests

$B$  bootstrap samples, tree  $T_b$  for each sample by recursively choosing the best out of  $d$  split points, return majority vote of trees.

## 10 Markov Decision Processes

$$\langle S, A, T, R, \gamma \rangle$$

↑      ↑  
states    actions

$$T: S \times A \times S \rightarrow \mathbb{R}$$

$$T(s, a, s') = \Pr(s_t = s' \mid s_{t-1} = s, a_{t-1} = a)$$

$R: S \times A \rightarrow \mathbb{R}$  reward for taking action  $a$  at state  $s$ .

$\gamma \in [0, 1]$  = discount factor.

policy  $\pi: S \rightarrow A$  specifies which action to take in each state

finite horizon value:  $V_\pi^h(s) = R(s, \pi(s)) + \gamma \sum_{s'} T(s, \pi(s), s') V_\pi^{h-1}(s')$

$\pi_1 \succ_h \pi_2$  if  $\forall s \in S, V_{\pi_1}^h(s) \geq V_{\pi_2}^h(s)$  and  $\exists s \in S$  s.t.  $V_{\pi_1}^h(s) > V_{\pi_2}^h(s)$ .  
 "better"

infinite horizon value: use  $\mathbb{E} \left[ \sum_{t=0}^{\infty} \gamma^t R_t \mid \pi, s_0 \right] = \mathbb{E} \left[ R_0 + \gamma R_1 + \dots \mid \pi, s_0 \right]$

$\uparrow$   
random variable  
representing reward  
across all possible  
stochastic transitions  
at time  $t$

### Bellman Equation

$$V_\pi(s) = \mathbb{E}[R_0 \mid \pi, s_0 = s] + \gamma \mathbb{E}[R_1 + \gamma(R_2 + \gamma(R_3 + \dots)) \mid \pi, s_0 = s]$$

$$= R(s, \pi(s)) + \gamma \sum_{s'} T(s, \pi(s), s') V_\pi(s')$$

finite policy finding:  $Q$  = "expected" value

$$Q^0(s, a) = 0$$

$$Q^1(s, a) = R(s, a) + 0$$

$$Q^2(s, a) = R(s, a) + \gamma \sum_{s'} T(s, a, s') \max_{a'} Q^1(s', a')$$

⋮

$$Q^h(s, a) = R(s, a) + \gamma \sum_{s'} T(s, a, s') \max_{a'} Q^{h-1}(s', a')$$

$$\pi_h^*(s, a) = \operatorname{argmax}_a Q^h(s, a)$$

infinite policy finding:

$\exists$  optimal policy  $\pi^*$  s.t.  $V_{\pi^*}(s) \geq V_\pi(s) \quad \forall s, \pi.$

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s'} T(s, a, s') \max_{a'} Q^*(s', a') \quad \pi^*(s, a) = \operatorname{argmax}_{a'} Q^*(s, a)$$

↑

expected infinite-horizon  
discounted value of being  
in state  $s$ , action  $a$ ,  
and doing  $\pi$  onwards

↳ update starting w/ initial  $Q(s, a) = 0$

↳ return  $Q$  when  $\|Q^{\text{new}}(s, a) - Q^{\text{old}}(s, a)\| < \epsilon \quad \forall s \in S, a \in A$

result is that we get  $\|V_{\pi_{Q^{\text{new}}}} - V_{\pi^*}\| < \epsilon$ . Also  $\exists$  some  $\epsilon' > 0$  s.t.  $\|Q_{\text{old}} - Q^{\text{new}}\| < \epsilon' \Rightarrow \pi_{Q^{\text{new}}} = \pi^*$

## II Reinforcement Learning

"Learning" an MDP without knowing  $R, T$ .

model-free method (Q-learning)

initialize all  $Q(s, a) = 0$ .  
given  $(s_{\text{old}}, a, s_{\text{new}}, r)$ , [pick  $a$  based on current  $\max_{a'} Q(s, a')$ ]

$$Q_{\text{new}}(s_{\text{old}}, a) = (1-\alpha) Q_{\text{old}}(s_{\text{old}}, a) + \alpha [r + \gamma \max_{a'} Q(s_{\text{new}}, a')]$$

to address exploitation vs. exploration,

we can also pick  $\arg \max_{a' \in A} Q(s, a')$  with  $1-\epsilon$ , random action with  $\epsilon$ .

deep Q-learning: minimize loss  $(Q(s, a) - (r + \gamma \max_{a'} Q(s, a')))^2$



using a neural network

catastrophic forgetting

← can handle using  
replay buffer

fitted Q-learning: alternate between running Q-function's policy  
and training neural network to get Q

policy gradient: define some form  $f$  differentiable that can  
be representative of  $\Pr[\text{action } a \mid \text{state } s]$ ,  
then gradient descent

## model-based RL

Solve the following MDP.

We know  $R(s, a)$  when we experience it

$$\text{set } \hat{T}(s, a, s') = \frac{\#(s, a, s') + 1}{\#(s, a) + |S|} \rightarrow \text{starts at } \frac{1}{|S|}$$

## (12) Unsupervised Learning

- no labels to learn.

clustering - given points, how to categorize in  $k$  groups

$$\min \left( \text{loss} = \sum_{j=1}^k \sum_{i=1}^n \mathbb{1}_{(y^{(i)}=j)} \|x^{(i)} - m_j^{(c)}\|^2 \right)$$

↑  
centroid  
of cluster

hyperparam.  
potential metrics  
for evaluation

- consistency
- ground truth (if it exists)
- visualization
- interpretability

$k$ -means : classify each point, by nearest neighbor  
then centroids at each new cluster,  
repeat

initial centroids /  $k$  both  
important choices.

autoencoders - encode vectors into shrunken dimension ( $k$ ), then decode  
(e.g. with linear/ReLU layers)

want small  $k$

can minimize squared loss

$$\min_{W^1, W_0^1, W^2, W_0^2} \sum_{i=1}^n \|h(g(x^{(i)}; W^1, W_0^1); W^2, W_0^2) - x^{(i)}\|^2$$

